

Homework 3

Please scan and upload your assignments on or before March 5, 2020.

- You are encouraged to discuss ideas with each other; but
 - you **must acknowledge** your collaborator, and
 - you **must compose your own** writeup and/or code independently.
 - We **strongly** encourage answers to theory questions in Latex, and answers to coding questions in Python (Jupyter notebooks).
 - Please upload your solutions in the form of a single .pdf or .zip file on NYUClasses.
 - Maximum score: 50 points.
-

1. (10 points) Suppose we wish to learn a regularized least squares model:

$$L(w) = \frac{1}{2} \sum_{i=1}^n (y_i - \langle w, x_i \rangle)^2 + \lambda R(w)$$

where $R(w)$ is a regularization function to be determined. Suggest good choices for $R(w)$ if the following criteria need to be achieved (there are no unique correct answers) and justify your choice in a sentence or two:

- a. All parameters w are free to be determined.
 - b. w should be sparse (i.e., only a few coefficients of w are nonzero).
 - c. The coefficients of w should be small in magnitude on average.
 - d. For most indices j , w_j should be equal to w_{j-1} .
 - e. w should have no negative-valued coefficients.
2. (10 points) The *Boston Housing Dataset* has been collected by the US Census Service and consists of 14 urban quality-of-life variables, with the last one being the median house price for a given town. Code for loading the dataset is provided at the end of this assignment. **Implement a linear regression model with ridge regression that predicts median house prices** from the other variables. Use 10-fold **cross validation** on 80-20 train-test splits and **report the final R^2 values** that you discovered. (You may want to preprocess your data to the range $[0, 1]$ in order to get meaningful results.)
3. (10 points) In class, we discussed the *lasso* objective, where the regularizer was chosen to be the ℓ_1 -norm. Here, we will derive an analytical closed form expression for the minimizer of a slightly simpler problem. Suppose x is a d -dimensional input and w is a d -dimensional variable. Show that the minimizer of the loss function:

$$L(w) = \frac{1}{2} \|x - w\|_2^2 + \lambda \|w\|_1$$

is given by:

$$w_i^* = \begin{cases} x_i - \lambda & \text{if } x_i > \lambda, \\ x_i + \lambda & \text{if } x_i < -\lambda, \\ 0 & \text{otherwise.} \end{cases}$$

4. **(20 points)** In this problem, we will implement **logistic regression trained with GD/SGD** and validate on synthetic training data.

- a. Suppose that the data dimension d equals 2. Generate two clusters of data points with **100 points each** (so that the total data size is $n = 200$), by sampling from Gaussian distributions centered at $(0.5, 0.5)$ and $(-0.5, -0.5)$. Call the data points x_i , and label them as $y_i = \pm 1$ depending on which cluster they originated from. Choose the variance of the Gaussian to be small enough so that the data points are sufficiently well separated. **Plot the data points on the 2D plane** to confirm that this is the case.
- b. (Derive your own GD routines; do *not* use sklearn functions here.) Train a logistic regression model that tries to minimize:

$$L(w) = - \sum_{i=1}^n y_i \log \frac{1}{1 + e^{-\langle w, x_i \rangle}} + (1 - y_i) \log \frac{e^{-\langle w, x_i \rangle}}{1 + e^{-\langle w, x_i \rangle}}$$

using **Gradient Descent (GD)**. Plot the decay of the training loss function as a function of number of iterations.

- c. Train the same logistic regression model, but this time using **Stochastic Gradient Descent (SGD)**. **Demonstrate that SGD exhibits a slower rate of convergence than GD, but is faster per-iteration**, and does not suffer in terms of final quality. You may have to play around a bit with the step-size parameters as well as mini-batch sizes to get reasonable answers.
 - d. Overlay the original plot of data points on the 2D data plane with the two (final) models that you obtained above in parts b and c to visualize correctness of your implementation.
5. **(optional)** How much time (in hours) did you spend working on this assignment?

```
import pandas as pd
import numpy as np
from sklearn.datasets import load_boston

boston_dataset = load_boston()
boston = pd.DataFrame(boston_dataset.data,
                      columns=boston_dataset.feature_names)
boston['MEDV'] = boston_dataset.target
boston.head()
```